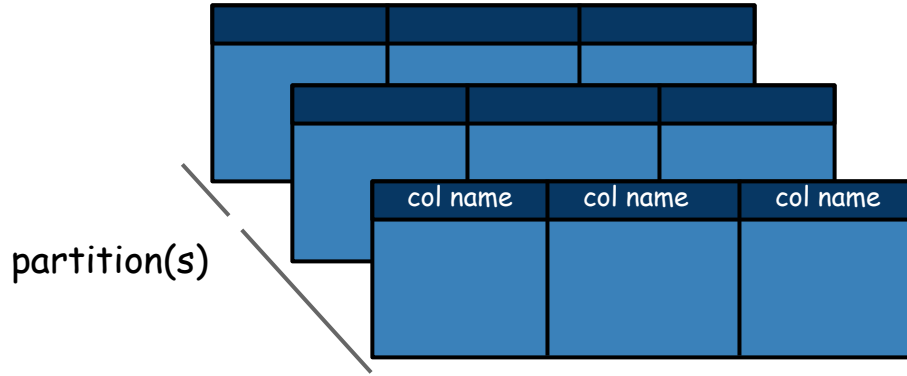



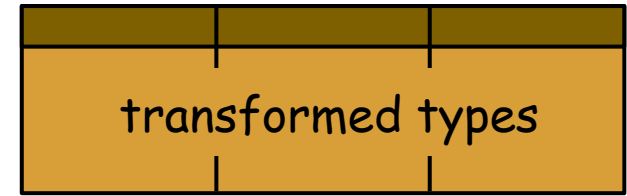
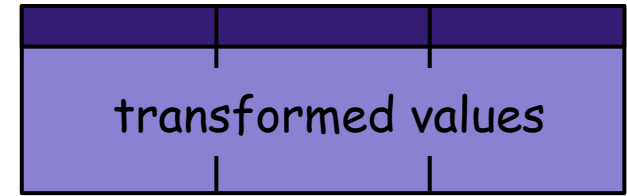
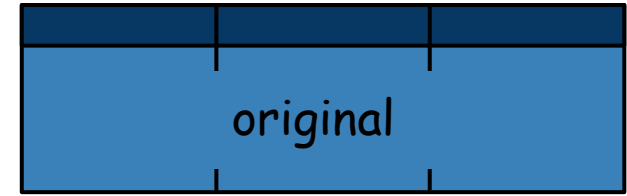
pyspark-pictures data frames

Learn the pyspark API through pictures and simple examples

data frame



row



user input

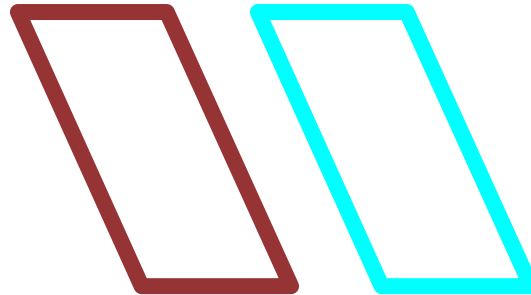
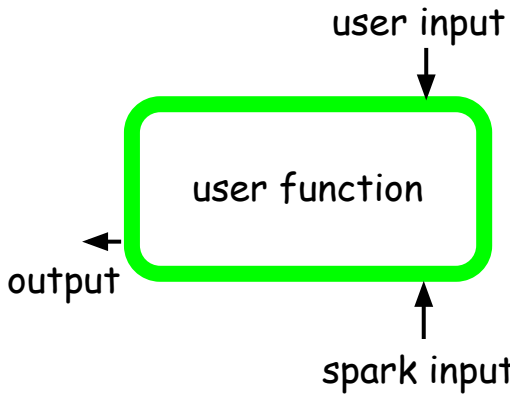
groupby
function

aggregate
function



user function

spark input

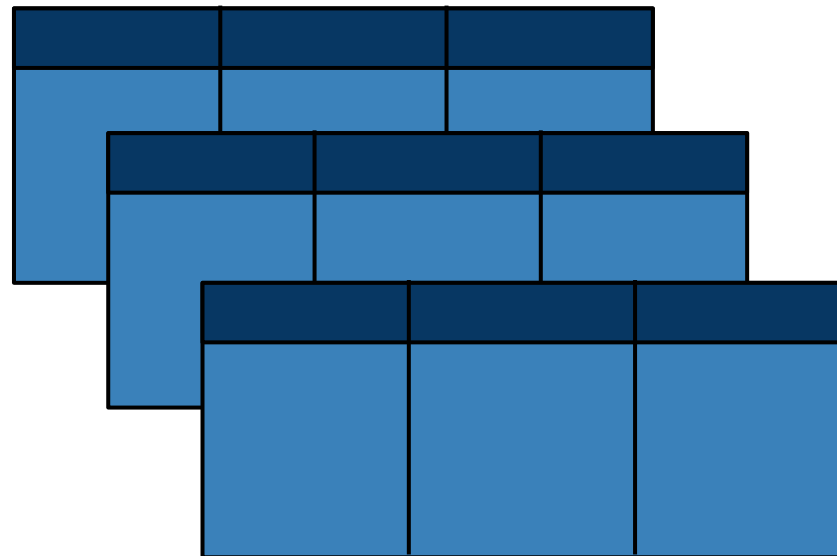
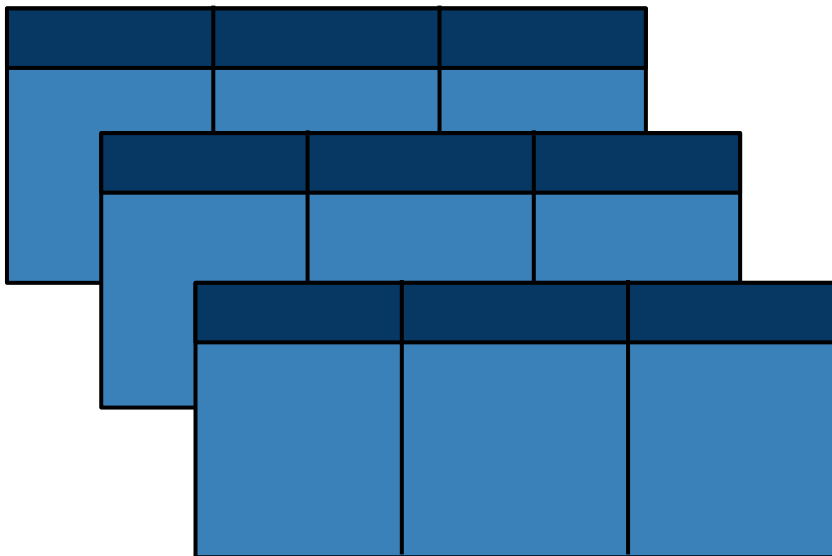
output



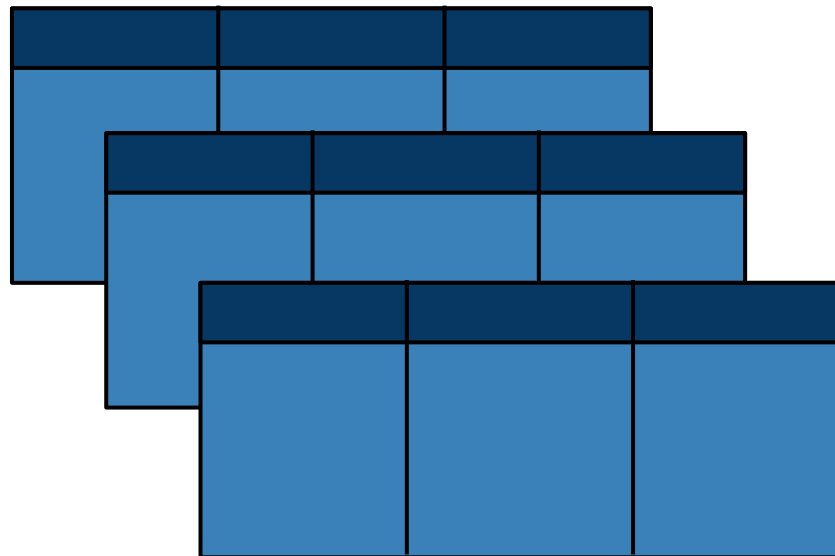
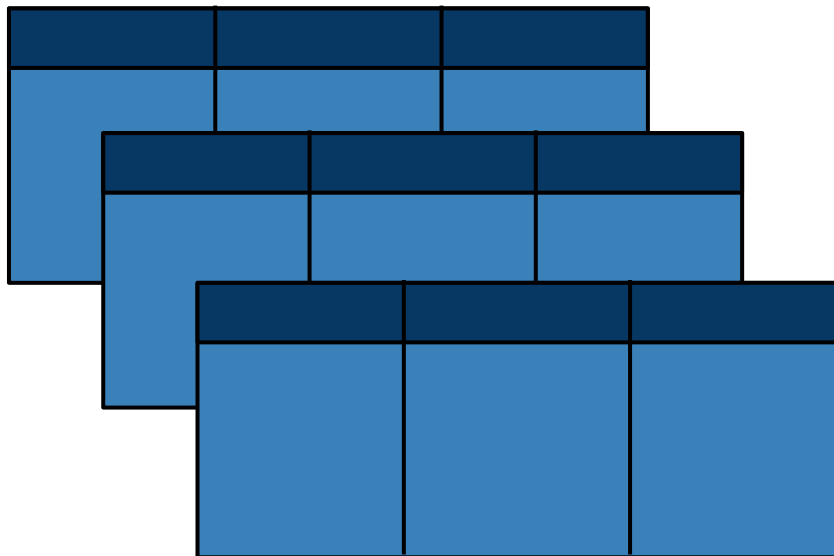
agg



alias

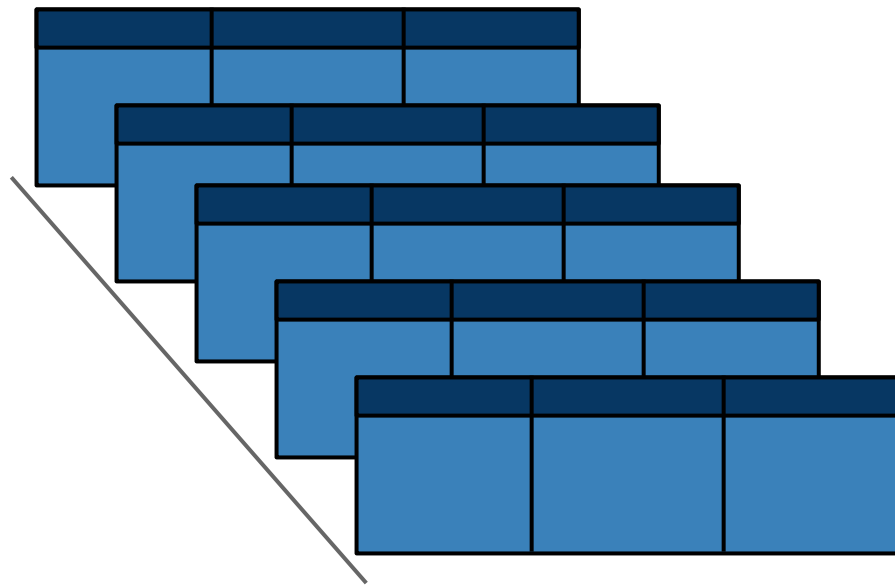
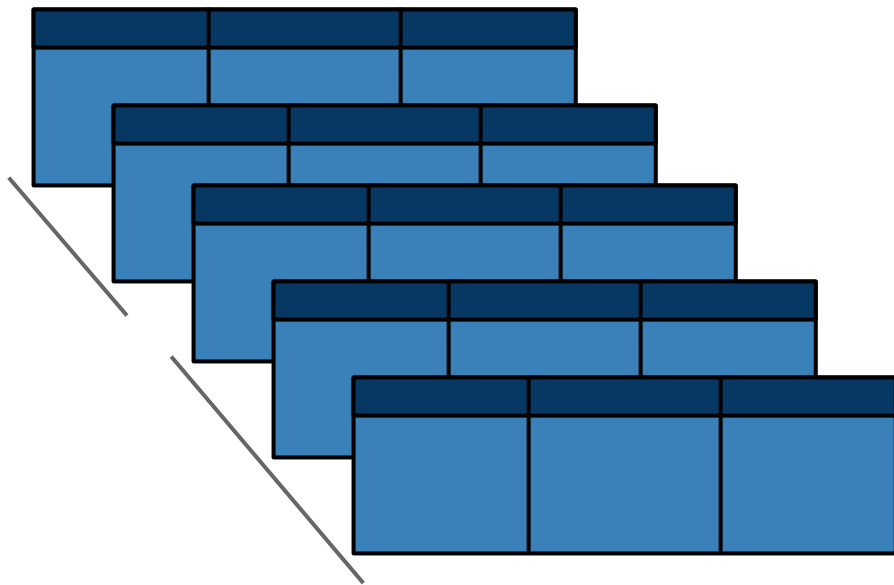


cache


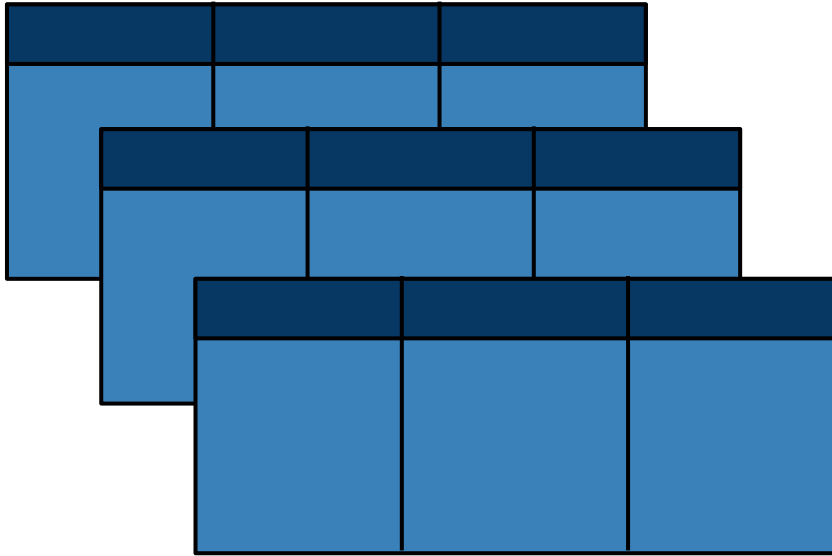


coalesce



numPartitions = 1




collect



columns



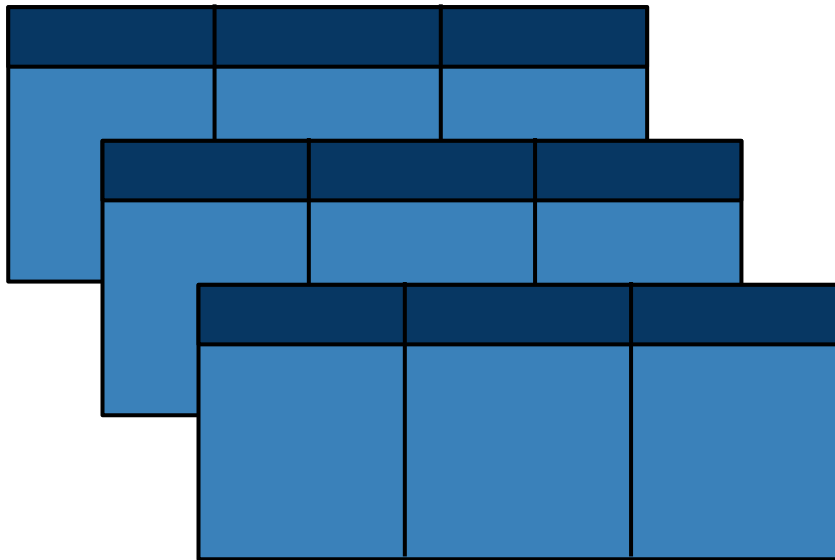
corr



Pearson's r

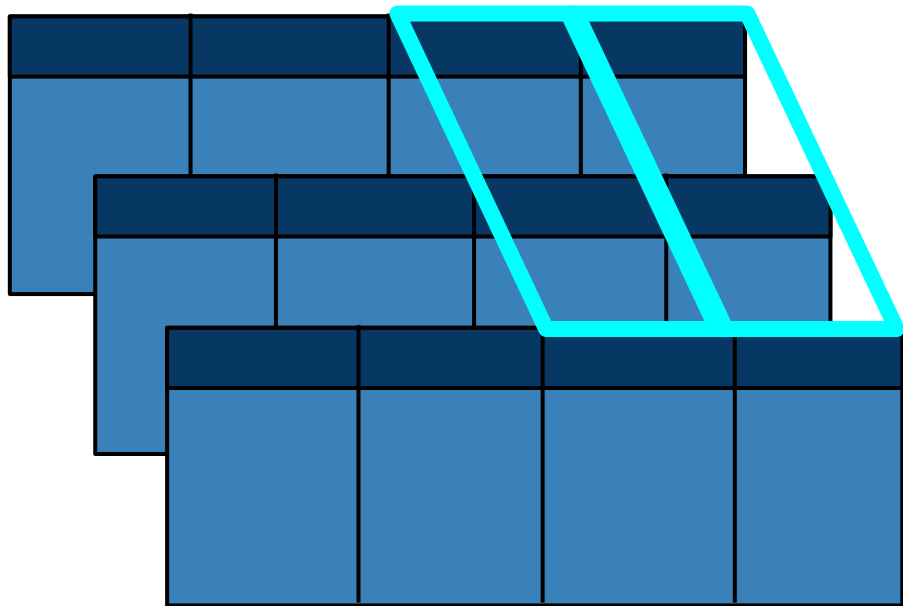
$$r = \frac{\sum_i (A_i - \bar{A})(C_i - \bar{C})}{\sqrt{\sum_i (A_i - \bar{A})^2} \sqrt{\sum_i (C_i - \bar{C})^2}}$$

count



3

COV

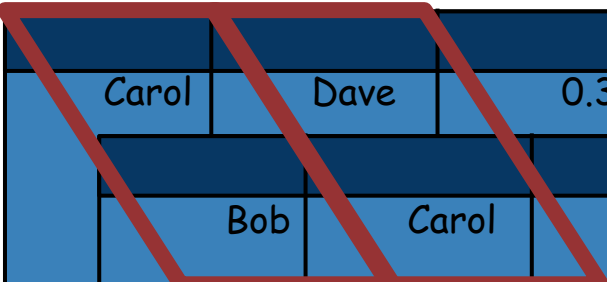


Sample Covariance

$$\frac{1}{N-1} \sum_i (A_i - \bar{A})(C_i - \bar{C})$$

crosstab

col1 = 'from' col2 = 'to'



Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

Carol	0	0	1
Bob	0	1	0
from_to	Bob	Carol	Dave
Alice	1	0	0

cube

*cols = 'from', 'to'

Alice	Carol	0.2
from	to	amt
Alice	Bob	0.1

null	null	
Alice	null	
null	Carol	
null	Bob	
Alice	Carol	
from	to	agg(amt)
Alice	Bob	

describe

Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

max	
min	
stdev	
mean	
summary	amt
count	

distinct

Bob	Carol	0.2
Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

drop

col = 'amt'

Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

Carol	Dave
Bob	Carol
from	to
Alice	Bob

dropDuplicates

subset = ['from', 'to']

Bob	Carol	0.2
Bob	Carol	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

dropna

how = 'any' subset = ['from', 'to']

Bob	Carol	0.2
Carol	null	0.3
Bob	Carol	null
from	to	amt
null	Bob	0.1

Bob	Carol	0.2
from	to	amt
Bob	Carol	null

dtypes

Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

[('from','string'), ('to','string'), ('amt','double')]

explain

extended = True

Carol	Dave	0.3

Bob	Carol	0.2

from	to	amt
Alice	Bob	0.1

== Parsed Logical Plan ==

...

== Analyzed Logical Plan ==

...

== Optimized Logical Plan ==

...

== Physical Plan ==

...

== RDD ==

fillna

value = 'unknown' subset = ['from', 'to']

Carol	null	0.3

Bob	Carol	nan

from	to	amt
null	Bob	0.1

Carol	unknown	0.3

Bob	Carol	nan

from	to	amt
unknown	Bob	0.1

filter

condition = "amt > 0.1"

Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

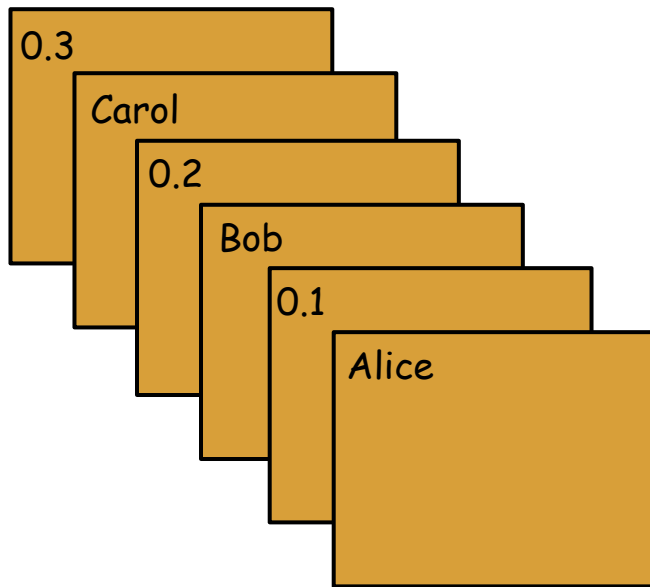
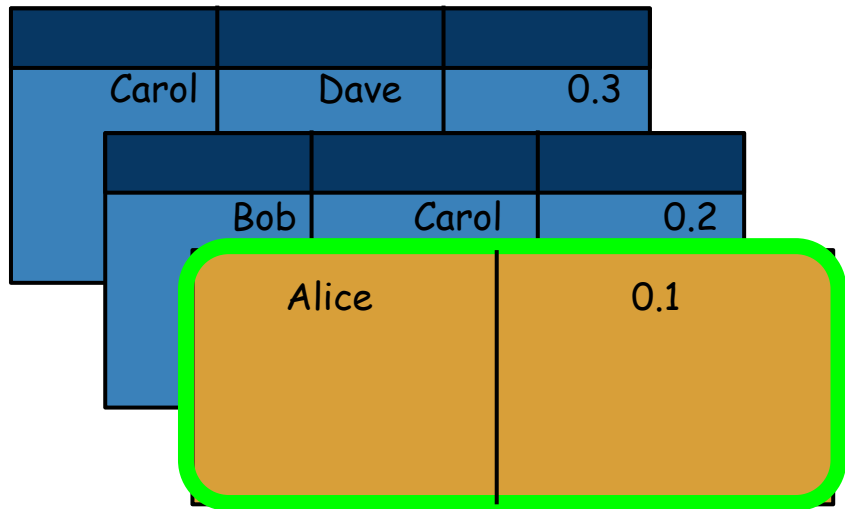
Carol	Dave	0.3
from	to	amt
Bob	Carol	0.2

first

	Carol	Dave 0.3
	Bob	Carol 0.2
	from	to amt
	Alice	Bob 0.1

Row(from='Alice', to='Bob', amt=0.1)

flatMap



foreach

Carol	Dave	0.3

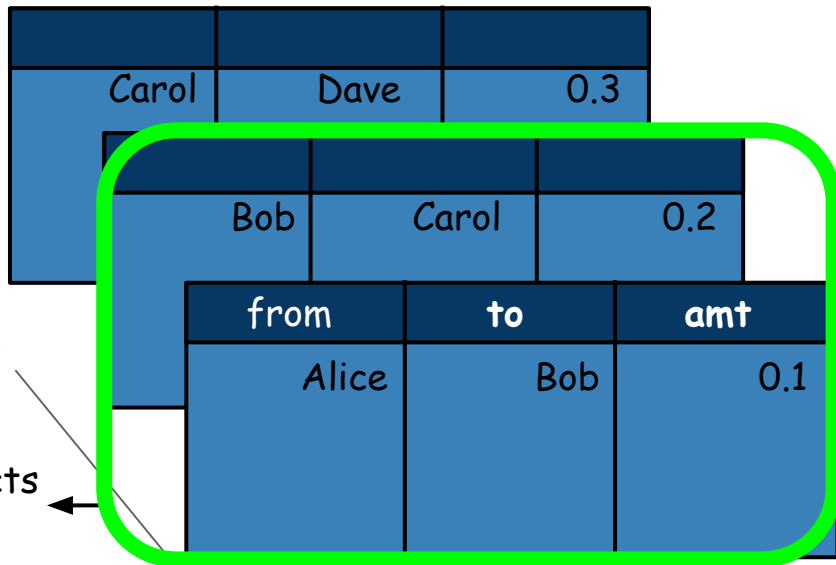
Bob	Carol	0.2

from	to	amt
Alice	Bob	0.1

side effects
(e.g print) ←

*no return value,
original DataFrame
unchanged

foreachPartition



The diagram illustrates the `foreachPartition` operation. It shows three nested DataFrames. The outermost DataFrame has a header row and one data row (Carol, Dave, 0.3). The middle DataFrame also has a header row and one data row (Bob, Carol, 0.2). The innermost DataFrame has a header row and one data row (Alice, Bob, 0.1). A green rounded rectangle highlights the innermost DataFrame, indicating it is the current partition being processed. An arrow points from the text 'side effects (e.g print)' to the green box.

Carol	Dave	0.3

Bob	Carol	0.2

from	to	amt
Alice	Bob	0.1

side effects
(e.g print)

*no return value,
original DataFrame
unchanged


freqItems

cols = ['from', 'amt'] support = 0.8

Carol	Bob	0.1
Alice	Bob	0.5
Alice	Bob	0.1
Alice	Dave	0.1
from	to	amt
Bob	Carol	0.1

from_freqItems	amt_freqItems
[Alice]	[0.1]

groupBy (groupby)



GroupedData Object
with methods: *agg, avg, count,*
max, mean, min, pivot, sum

groupBy(col1).avg(col2)

col1 = 'from' col2 = 'amt'

	Carol	Dave	0.3
	Alice	Carol	0.2
from	to	amt	
Alice	Bob	0.1	

Carol	0.3
from	avg(amt)
Alice	0.15

head

n = 2

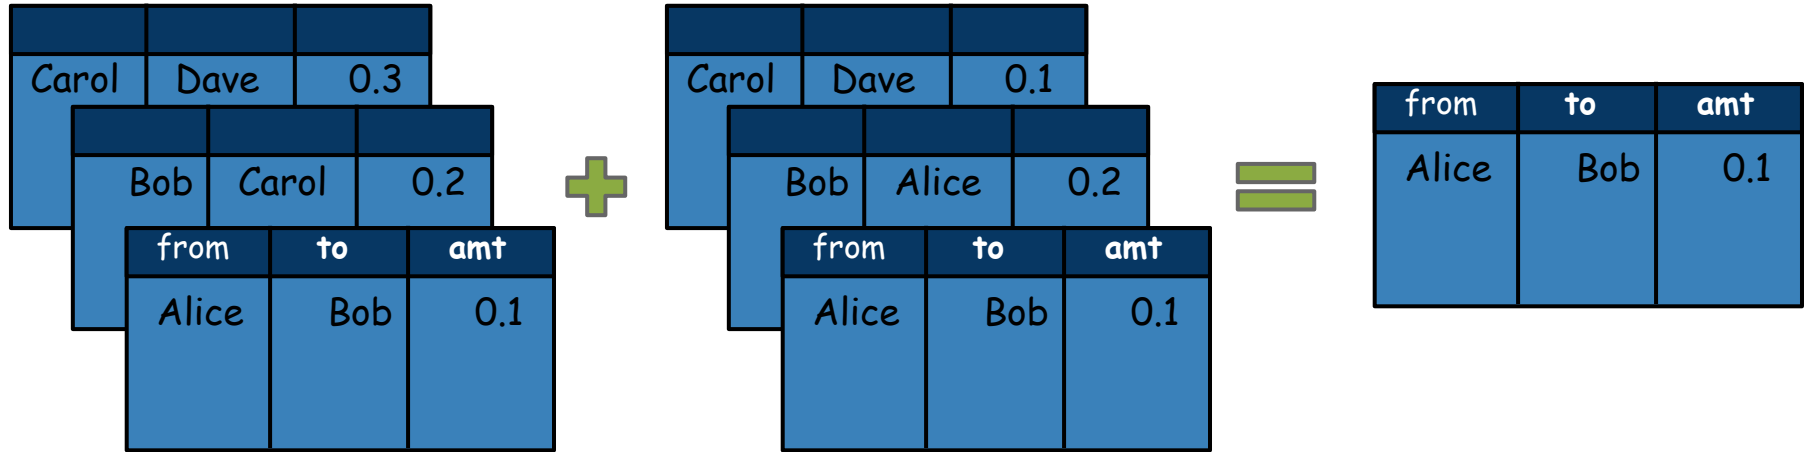
Carol	Dave	0.3

Bob	Carol	0.2


from	to	amt
Alice	Bob	0.1

[Row(from=u'Alice', to=u'Bob', amt=0.1),
Row(from=u'Bob', to=u'Carol', amt=0.2)]

intersect



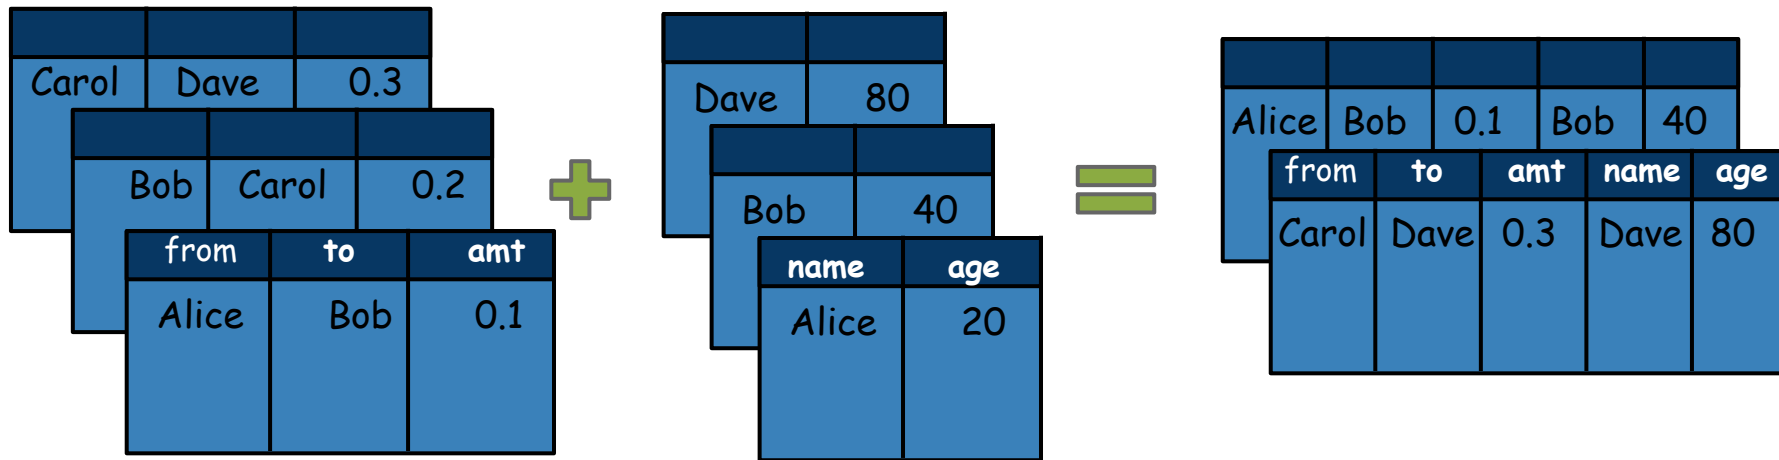
isLocal



False

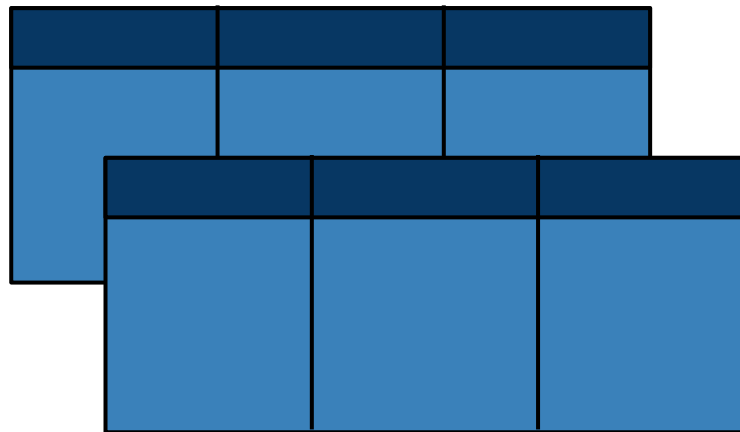
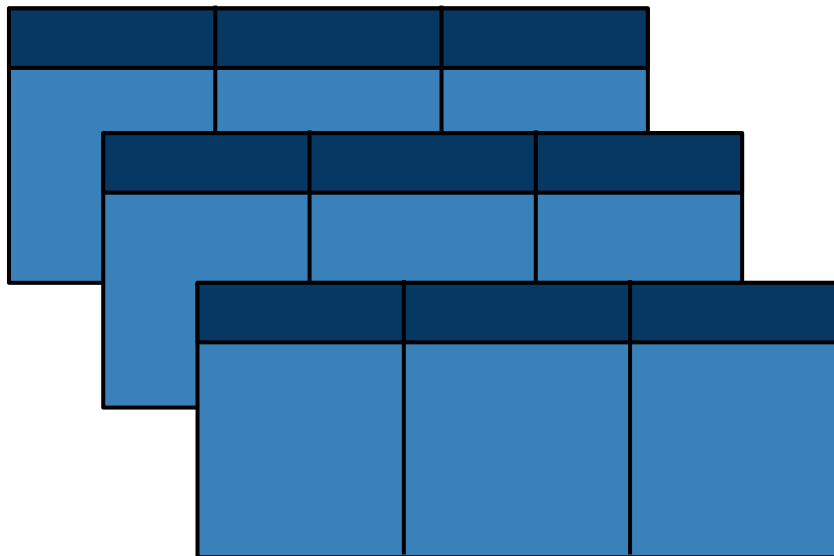
join

joinExprs = x.to==y.name joinType = 'inner'




limit

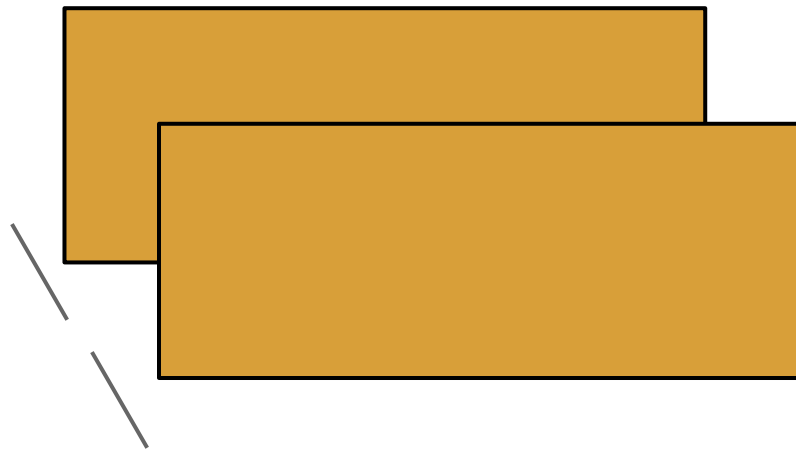

num = 2



map



mapPartitions



na

Bob	Carol	0.2

Carol	null	0.3

Bob	Carol	null

from	to	amt
null	Bob	0.1

DataFrameNaFunctions
Object

with methods: drop, fill, replace

orderBy



```
cols = ['from'], ascending = [False]
```

Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

[illegible]

persist

```
storageLevel =  
StorageLevel(MEMORY_ONLY_SER)
```



printSchema

Carol	Dave	0.3

Bob	Carol	0.2

from	to	amt
Alice	Bob	0.1

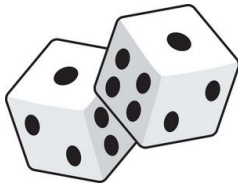
stdout

root

|-- from: string (nullable = true)
|-- to: string (nullable = true)
|-- amt: double (nullable = true)

randomSplit

weights = [0.5,0.5]

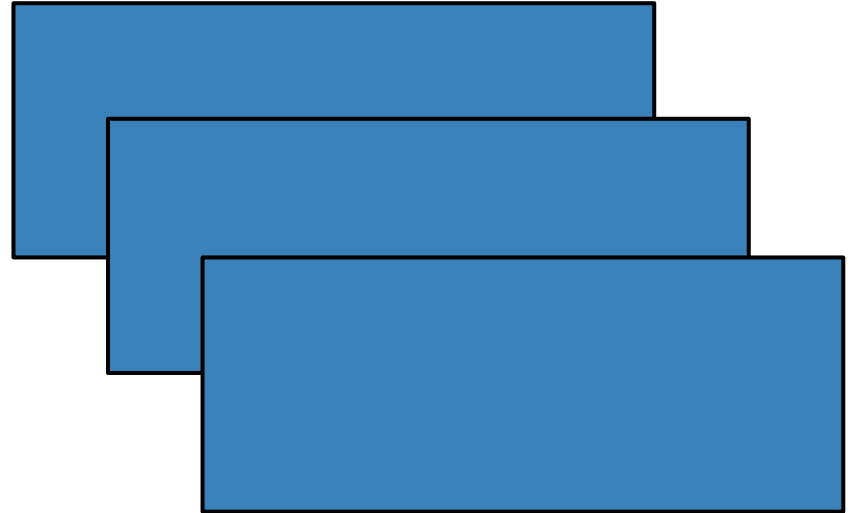
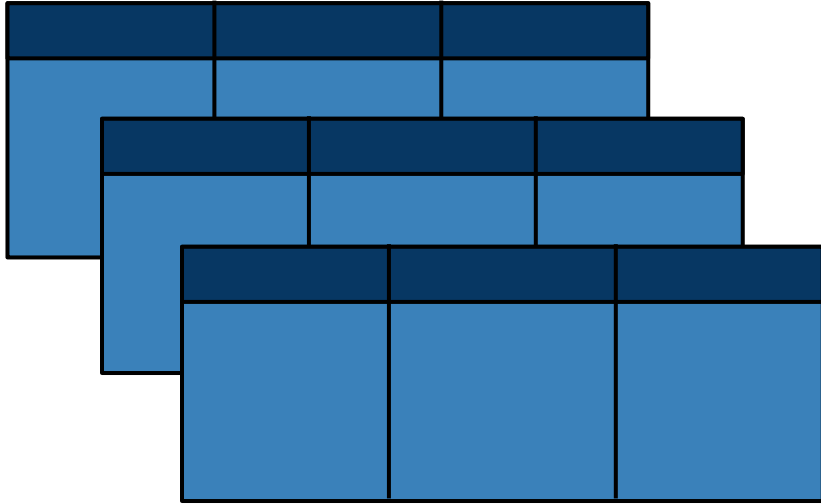


Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

Carol	Dave	0.3
from	to	amt
Bob	Carol	0.2

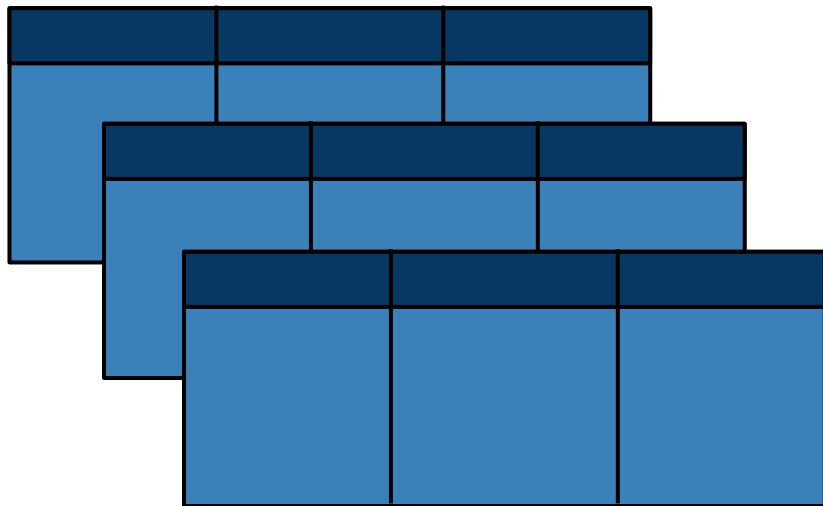
from	to	amt
Alice	Bob	0.1

rdd



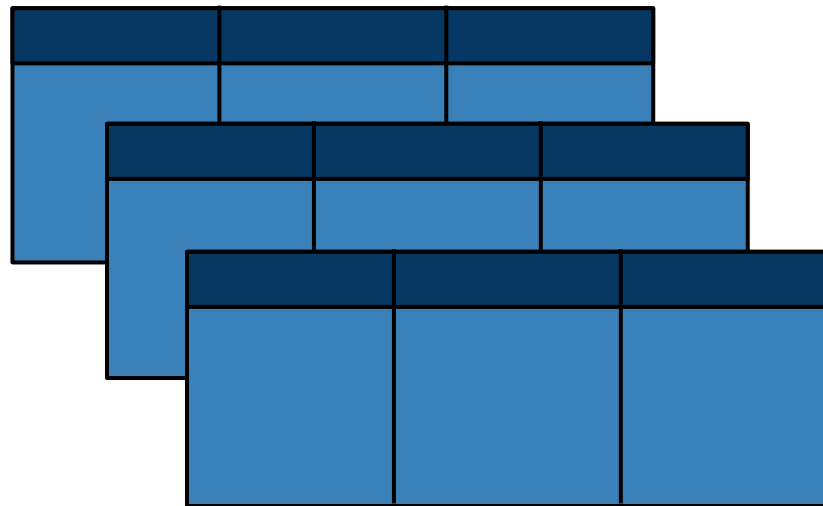
registerTempTable

name = "TRANSACTIONS"



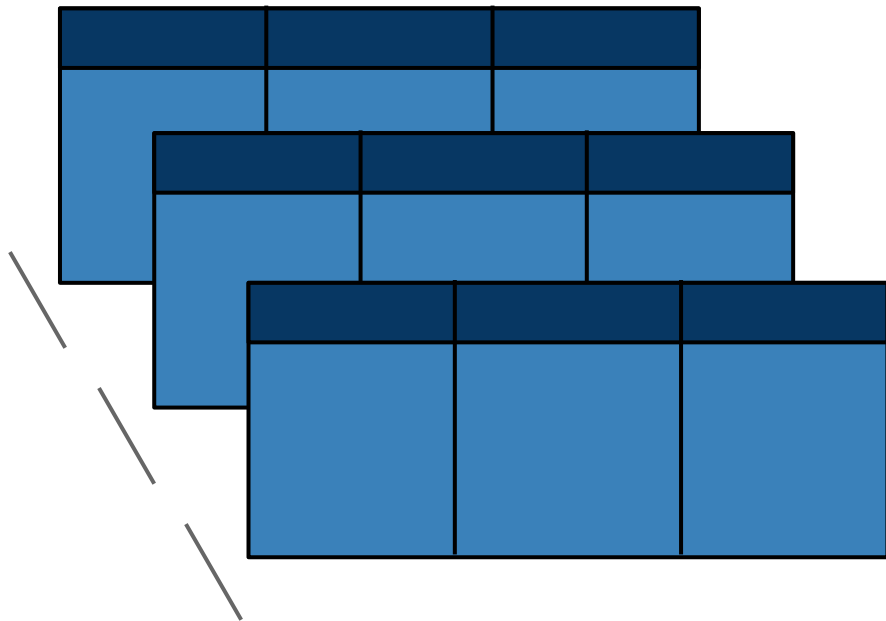
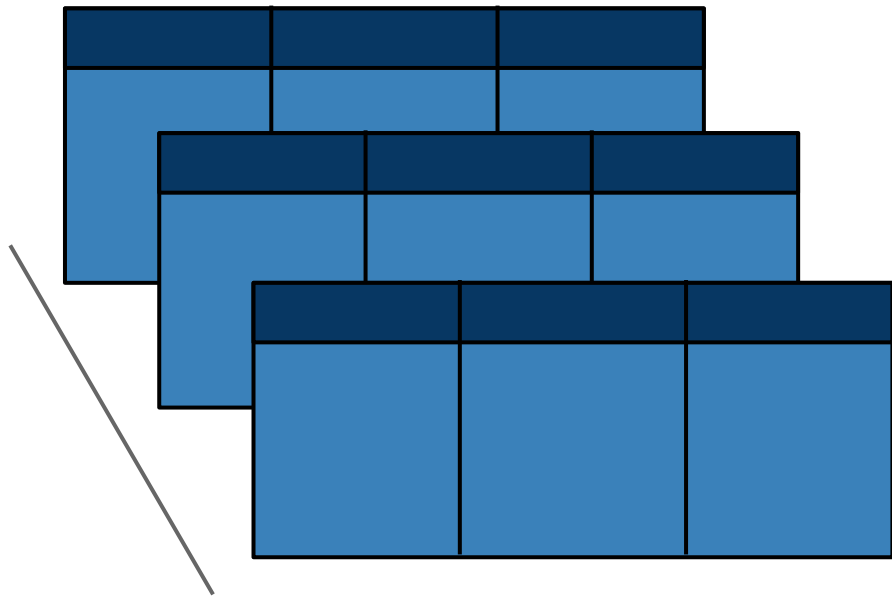
The diagram shows three nested table blocks, each with a dark blue header and light blue body. The top block is 3 columns wide, the middle is 4 columns wide, and the bottom is 5 columns wide. They are arranged in a staggered, overlapping fashion from top-left to bottom-right.

TRANSACTIONS



The diagram shows three nested table blocks, each with a dark blue header and light blue body. The top block is 3 columns wide, the middle is 4 columns wide, and the bottom is 5 columns wide. They are arranged in a staggered, overlapping fashion from top-left to bottom-right, identical to the temporary table structure.

repartition



replace

```
to_replace = 'Dave' value = 'David'
```

Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

Carol	David	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

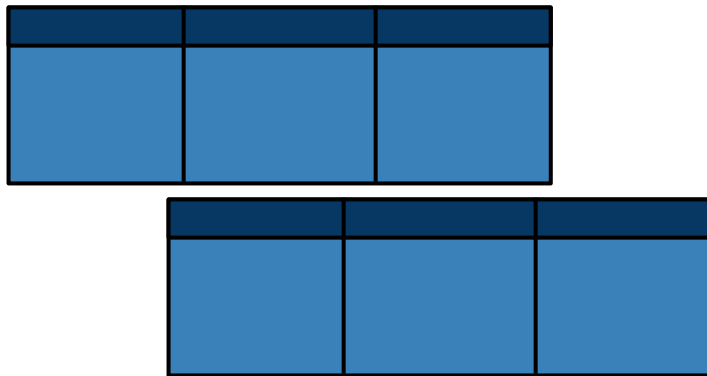
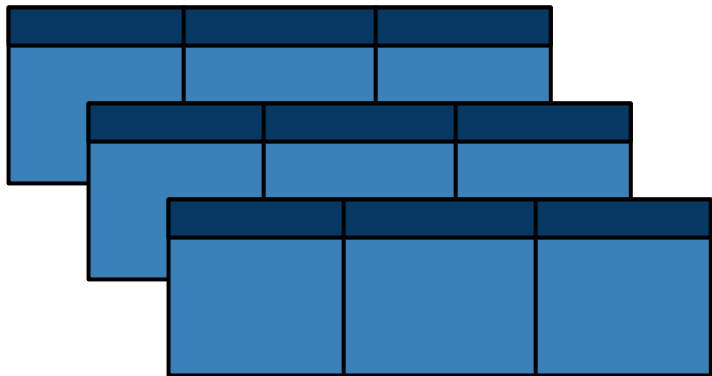
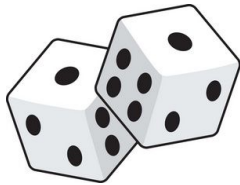
rollup

cols = ['from', 'to']

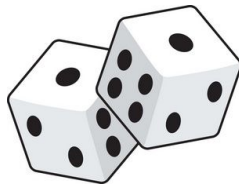
Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

null	null	
Carol	null	
Bob	null	
Carol	Dave	
Alice	null	
Bob	Carol	
from	to	agg(amt)
Alice	Bob	

sample



sampleBy



Bob	Carol	0.6
Bob	Bob	0.5
Alice	Dave	0.4
Alice	Alice	0.3
Alice	Carol	0.2
from	to	amt
Alice	Bob	0.1

Bob	Carol	0.6
Bob	Bob	0.5
from	to	amt
Alice	Bob	0.3

schema

from	to	amt

from	to	amt
------	----	-----

select

```
cols = ['from', 'amt']
```

from	to	amt

from	amt

selectExpr

```
expr = ["substr(from,1,1)", "amt + 10"]
```

Carol	Dave	0.3

Bob	Carol	0.2

from	to	amt
Alice	Bob	0.1

The diagram illustrates three overlapping tables. The top table has columns *C* and 10.3. The middle table has columns *B* and 10.2. The bottom table has columns *from* and *amt*, with values *A* and 10.1 respectively.

show

Carol	Dave	0.3

Bob	Carol	0.2

from	to	amt
Alice	Bob	0.1

stdout

```
+-----+-----+---+  
| from|  to|amt|  
+-----+-----+---+  
| Alice| Bob|0.1|  
| Bob|Carol|0.2|  
| Carol| Dave|0.3|  
+-----+-----+---+
```

sort

cols = ['to']

Carol	Alice	0.3	
Bob	Carol	0.2	
	from	to	amt
	Alice	Bob	0.1


Bob	Carol	0.2

sortWithinPartitions

Carol	Alice	0.3	2
Bob	Carol	0.2	2
from	to	amt	p_id
Alice	Bob	0.1	1

Bob	Carol	0.2	2
Carol	Alice	0.3	2
from	to	amt	p_id
Alice	Bob	0.1	1

stat



DataFrameStatFunctions
Object
with methods: `corr`, `cov`,
`corsstab`, `freqItems`, `sampleBy`

subtract

from	to	amt
Carol	Dave	0.3
Bob	Carol	0.2
Alice	Bob	0.1

The diagram illustrates a graph structure with three overlapping tables. The top table has columns Carol, Dave, and 0.1. The middle table has columns Bob, Carol, and 0.2. The bottom table has columns from, to, and amt, with rows Alice to Bob (0.1) and Bob to Carol (0.2). A green arrow points from the bottom table to the middle table.

Carol	Dave	0.1


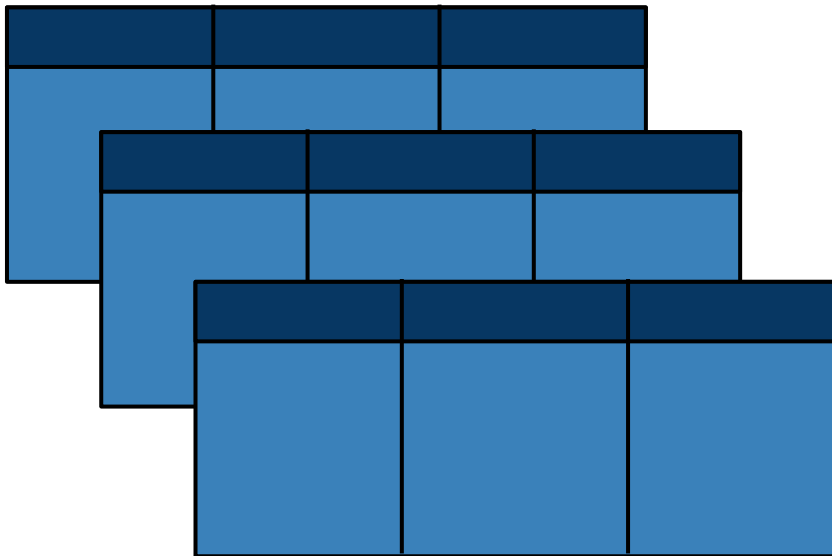
Bob	Carol	0.2

from	to	amt
Alice	Bob	0.1
	Bob	0.2

from	to	amt
Carol	Dave	0.3

take

num = 2



toDF

```
cols = ["seller", "buyer"]
```

from	to	amt

seller	buyer	amt

toJSON

Carol	Alice	0.3

Bob	Carol	0.2

from	to	amt
Alice	Bob	0.1

u'{"from":"Carol","to":"Alice","amt":0.3}'

u'{"from":"Bob","to":"Carol","amt":0.2}'

u'{"from":"Alice","to":"Bob","amt":0.1}'

toPandas

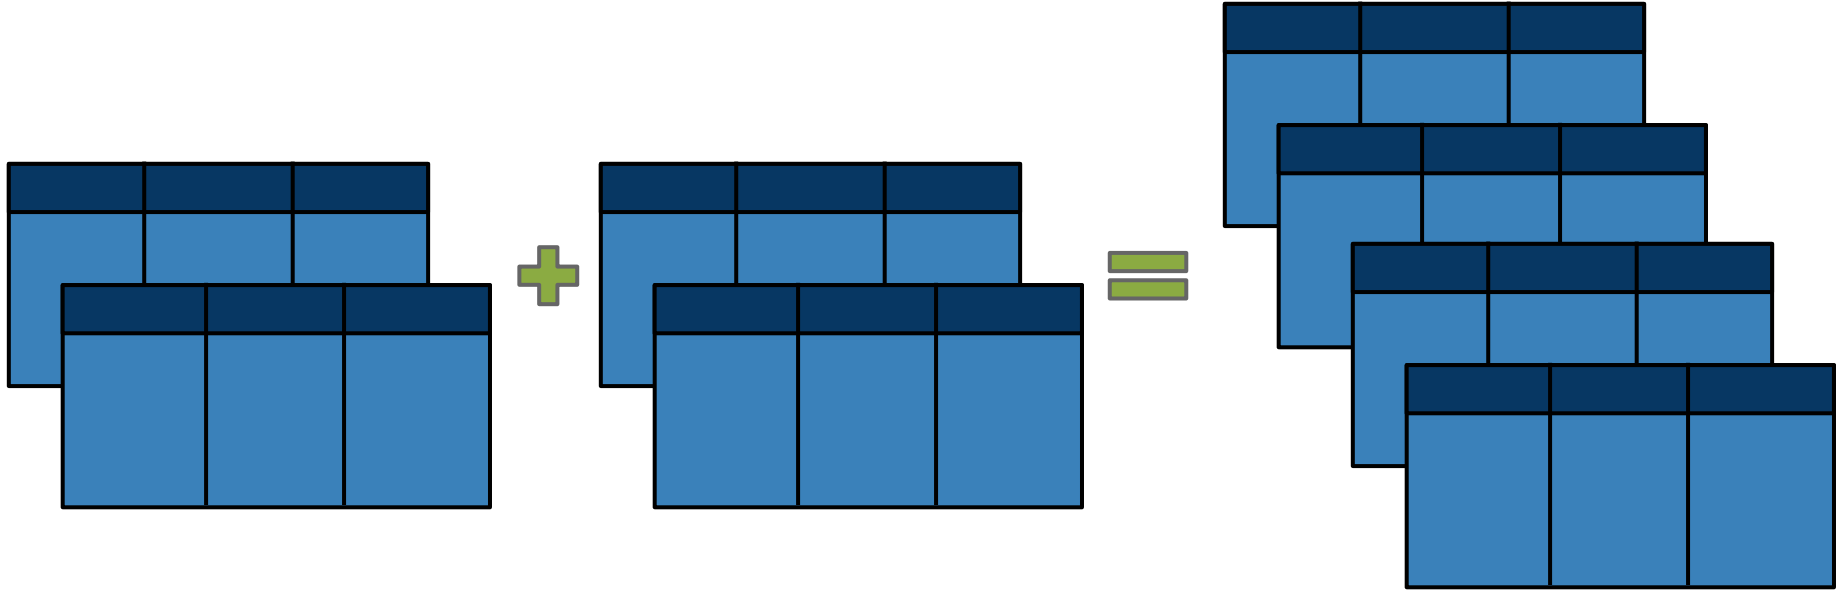
Carol	Alice	0.3

Bob	Carol	0.2


from	to	amt
Alice	Bob	0.1

```
from to  amt
0 Alice Bob 0.1
1 Bob  Carol 0.2
2 Carol Alice 0.3
```

unionAll



unpersist



where (filter)

condition = "amt > 0.1"

Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

Carol	Dave	0.3
from	to	amt
Bob	Carol	0.2

withColumn

colName = 'conf'

from	to	amt

from	to	amt	conf

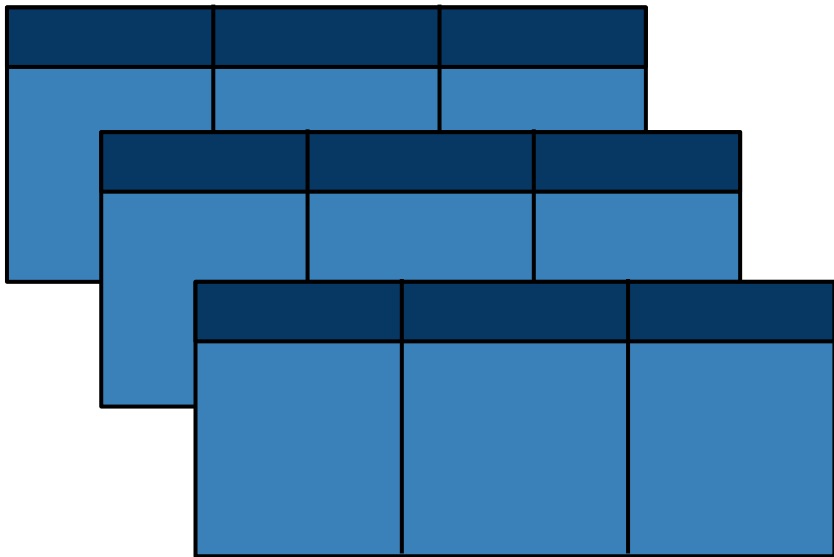
withColumnRenamed

```
existing = 'amt' col = 'amount'
```

from	to	amt

from	to	amount

write



DataFrameWriter
Object

with methods: format,
insertInto, jdbc, json, mode,
option, options, orc, parquet,
partitionBy, save, saveAsTable,
text